

Single-Source Shortest Paths

Shortest Path Problems

- How can we find the shortest route between two points on a map?
- Model the problem as a graph problem:
 - Road map is a weighted graph:
 - vertices** = cities
 - edges** = road segments between cities
 - edge weights** = road distances
 - Goal: find a shortest path between two vertices (cities)

Many applications


- Shortest paths model many useful real-world problems.
 - Minimization of latency in the Internet.
 - Minimization of cost in power delivery.
 - Job and resource scheduling.
 - Route planning.

Paths in graphs

- Consider a directed graph $G = (V, E)$ with edge-weight function $w : E \rightarrow \mathbb{R}$. The **weight** of path $p = v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_k$ is defined to be

$$w(p) = \sum_{i=1}^k w(v_{i-1}, v_i)$$

Example:



Shortest paths

- A *shortest path* from u to v is a path of minimum weight from u to v . The *shortest-path weight* from u to v is defined as

$$d(u, v) = \begin{cases} \min\{w(p) : p \text{ is a path from } u \text{ to } v\} \\ \infty \text{ if no path from } u \text{ to } v \text{ exists} \end{cases}$$

Optimal substructure

Lemma 24.1: Subpaths of shortest paths are shortest paths

Given a weighted, directed graph $G = (V, E)$ with weight function $w: E \rightarrow R$, let $p = (v_0, v_1, \dots, v_k)$ be a shortest path from vertex v_0 to vertex v_k and for any i and j such that $0 \leq i \leq j \leq k$, let $p_{ij} = (v_i, v_{i+1}, \dots, v_j)$ be the subpath of p from vertex i to vertex j . Then, p_{ij} is a shortest path from i to j .


Proof: If we decompose path p into $v_0 \xrightarrow{p_{0i}} v_i \xrightarrow{p_{ij}} v_j \xrightarrow{p_{jk}} v_k$, then we have that $w(p) = w(p_{0i}) + w(p_{ij}) + w(p_{jk})$.

Assume that there is a path p'_{ij} from i to j with weight $w(p'_{ij}) < w(p_{ij})$. Then, $v_0 \xrightarrow{p_{0i}} v_i \xrightarrow{p'_{ij}} v_j \xrightarrow{p_{jk}} v_k$ is a path from 0 to k whose weight $w(p_{0i}) + w(p'_{ij}) + w(p_{jk})$ is less than $w(p)$, which contradicts the assumption that p is a shortest path from 0 to k .

Shortest-Path Representation

For each vertex $v \in V$:

- $d[v] = \delta(s, v)$: a **shortest-path estimate**
 - Initially, $d[v] = \infty$
 - Reduces as algorithms progress
- $\pi[v] =$ **predecessor** of v on a shortest path from s
 - If no predecessor, $\pi[v] = \text{NIL}$
 - π induces a tree—**shortest-path tree**
- Shortest paths & shortest path trees are not unique



Initialization

INITIALIZE-SINGLE-SOURCE(V, s)

1. **for** each $v \in V$
2. **do** $d[v] \leftarrow \infty$
3. $\pi[v] \leftarrow \text{NIL}$
4. $d[s] \leftarrow 0$

All the shortest-paths algorithms start with
INITIALIZE-SINGLE-SOURCE



Relaxation

- **Relaxing** an edge (u, v) = testing whether we can improve the shortest path to v found so far by going through u

If $d[v] > d[u] + w(u, v)$

we can improve the shortest path to v

\Rightarrow update $d[v]$ and $\pi[v]$



After relaxation:

$$d[v] \leq d[u] + w(u, v)$$

RELAX(u, v, w)

1. if $d[v] > d[u] + w(u, v)$
 2. then $d[v] \leftarrow d[u] + w(u, v)$
 3. $\pi[v] \leftarrow u$
- All the single-source shortest-paths algorithms
 - start by calling INIT-SINGLE-SOURCE
 - then relax edges
 - The algorithms differ in the order and how many times they relax each edge





Dijkstra's Algorithm

- Single-source shortest path problem:
 - No negative-weight edges: $w(u, v) > 0 \forall (u, v) \in E$
- Maintains two sets of vertices:
 - S = vertices whose final shortest-path weights have already been determined
 - Q = vertices in $V - S$: min-priority queue
 - Keys in Q are estimates of shortest-path weights ($d[v]$)
- Repeatedly select a vertex $u \in V - S$, with the minimum shortest-path estimate $d[v]$



Dijkstra (G, w, s)

1. INITIALIZE-SINGLE-SOURCE(V, s)
2. $S \leftarrow \emptyset$
3. $Q \leftarrow V[G]$
4. **while** $Q \neq \emptyset$
5. $u \leftarrow \text{EXTRACT-MIN}(Q)$
6. $S \leftarrow S \cup \{u\}$
7. **for** each vertex $v \in \text{Adj}[u]$
8. RELAX(u, v, w)

Dijkstra's Algorithm(Example)



Dijkstra's Algorithm(Example)



Dijkstra (G, w, s)

1. INITIALIZE-SINGLE-SOURCE(V, s) $\leftarrow \Theta(V)$
2. $S \leftarrow \emptyset$
3. $Q \leftarrow V[G]$ $\leftarrow O(V)$ build min-heap
4. **while** $Q \neq \emptyset$ $\leftarrow O(V)$
5. $u \leftarrow \text{EXTRACT-MIN}(Q)$ $\leftarrow O(\lg V)$
6. $S \leftarrow S \cup \{u\}$
7. **for** each vertex $v \in \text{Adj}[u]$ $\leftarrow O(E)$
8. RELAX(u, v, w) $\leftarrow O(\lg V)$

Running time: $O(V \lg V + E \lg V) = O(E \lg V)$

Analysis of Dijkstra's Algorithm

$$\text{Time} = \Theta(n) \cdot T_{\text{EXTRACTMIN}} + \Theta(m) \cdot T_{\text{ChangeKEY}}$$


Q	$T_{\text{EXTRACTMIN}}$	$T_{\text{ChangeKEY}}$	Total
array	$\Theta(n)$	$\Theta(1)$	$\Theta(n^2)$
Priority queue	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(m \log n)$

Correctness (Dijkstra's Algorithm)


Prove in each iteration, $u.d = \delta(s, u)$ for the vertex added to set $S(\text{visited})$.

Let u be the first vertex for which $u.d \neq \delta(s, u)$ when it is added to set S .

Let us consider the first vertex y along $\mathbf{p}(s \xrightarrow{p_1} x \rightarrow y \xrightarrow{p_2} u)$ such that $y \in V - S$, and let $x \in S$ be y 's predecessor along \mathbf{p} .



Correctness (Dijkstra's Algorithm)



We claim that $y.d = \delta(s, y)$ when u is added to S . Because we had $x.d = \delta(s, x)$ when x is added to S . $\text{Edge}(x, y)$ was relaxed at that time.

As y appears before u on a shortest path from s to u , we have $\delta(s, y) \leq \delta(s, u)$

$$y.d = \delta(s, y)$$

$$\leq \delta(s, u)$$

$$\leq u.d$$

But because both vertices u and y were in $V-S$

when u was chosen by extracting min from the

heap we have $u.d \leq y.d$


Therefore, $y.d = \delta(s, y) = \delta(s, u) = u.d$ which contradicts our choice of u .

Dijkstra's Algorithm - negative weights?

Dijkstra's Algorithm fails if there are negative weights.

Example: Select vertex v immediately after s

But shortest path from s to v is $s-x-y-v$



Bellman-Ford Idea


- Consider each edge (u,v) and see if u offers v a cheaper path from s
 - compare $d[v]$ to $d[u] + w(u,v)$
- Repeat this process $|V| - 1$ times to ensure that accurate information propagates from s , no matter what order the edges are considered in

Bellman-Ford Algorithm





- Single-source shortest paths problem
 - Computes $d[v]$ and $\pi[v]$ for all $v \in V$
- Allows negative edge weights
- Returns:
 - **TRUE** if no negative-weight cycles are reachable from the source s
 - **FALSE** otherwise \Rightarrow no solution exists
- Idea:
 - Traverse all the edges $|V - 1|$ times, every time performing a relaxation step of each edge

Bellman-Ford(G, w, s)

1. INITIALIZE-SINGLE-SOURCE(G, s)
2. **for** $i \leftarrow 1$ to $|G.V| - 1$
3. **do for** each edge $(u, v) \in G.E$
4. **do** RELAX(u, v, w)
5. **for** each edge $(u, v) \in G.E$
6. **do if** $d[v] > d[u] + w(u, v)$
7. **then return** FALSE
8. **return** TRUE






Bellman-Ford(Example)



Detecting Negative Cycles

for each edge $(u, v) \in E$
 do if $d[v] > d[u] + w(u, v)$
 then return FALSE
return TRUE



Look at edge (s, b) :

$$d[b] = -1$$

$$d[s] + w(s, b) = -4$$

$$\Rightarrow d[b] > d[s] + w(s, b)$$

Bellman-Ford(G, w, s)

- | | | | |
|----|--|------------------------|-----------|
| 1. | INITIALIZE-SINGLE-SOURCE(G, s) | $\leftarrow \Theta(V)$ | } $O(VE)$ |
| 2. | for $i \leftarrow 1$ to $ G.V - 1$ | $\leftarrow O(V)$ | |
| 3. | do for each edge $(u, v) \in G.E$ | $\leftarrow O(E)$ | |
| 4. | do RELAX(u, v, w) | | |
| 5. | for each edge $(u, v) \in G.E$ | $\leftarrow O(E)$ | |
| 6. | do if $d[v] > d[u] + w(u, v)$ | | |
| 7. | then return FALSE | | |
| 8. | return TRUE | | |

Running time: $O(VE)$

Correctness (Bellman-Ford)

If G does contain a negative-weight cycle reachable from s , then the algorithm returns FALSE.

Graph $G=(V,E)$ contains a negative-weight cycle $c = \langle v_0, v_1, \dots, v_k \rangle$ reachable from the source vertex s where $v_0 = v_k$. Then,

$$\sum_{i=1}^k w(v_{i-1}, v_i) < 0$$

Assume for the purpose of contradiction that the Bellman-Ford algorithm returns TRUE.

Thus, $v_i.d \leq v_{i-1}.d + w(v_{i-1}, v_i)$ for $i = 1, 2, \dots, k$

Correctness (Bellman-Ford)

$$\begin{aligned}\sum_{i=1}^k v_i \cdot d &\leq \sum_{i=1}^k (v_{i-1} \cdot d + w(v_{i-1}, v_i)) \\ &= \sum_{i=1}^k (v_{i-1} \cdot d) + \sum_{i=1}^k w(v_{i-1}, v_i)\end{aligned}$$

Since $v_0 = v_k$, each vertex in c appears exactly once in each of the summations.

$$\begin{aligned}\sum_{i=1}^k v_i \cdot d &= \sum_{i=1}^k v_{i-1} \cdot d \\ &+ \sum_{i=1}^k w(v_{i-1}, v_i) > 0\end{aligned}$$

which contradicts previous inequality.